

Import Required Libraries

```
In [1]: import pandas as pd
import numpy as np
```

Load Data

```
In [2]: movies_credits = pd.read_csv('./data/tmdb_5000_credits.csv')

movies = pd.read_csv('./data/tmdb_5000_movies.csv')
```

1.0 Initial Data Analysis

```
In [3]: # Top 10 movies ranked by the number of votes

movies.sort_values('vote_count', ascending=False).head(10)['title']
```

```
Out[3]: 96          Inception
65      The Dark Knight
0        Avatar
16      The Avengers
788      Deadpool
95      Interstellar
287      Django Unchained
94      Guardians of the Galaxy
426      The Hunger Games
127      Mad Max: Fury Road
Name: title, dtype: object
```

```
In [4]: # Top 10 movies ranked by the revenue

movies.sort_values('revenue', ascending=False).head(10)['title']
```

```
Out[4]: 0          Avatar
25      Titanic
16      The Avengers
28      Jurassic World
44      Furious 7
7      Avengers: Age of Ultron
124     Frozen
31      Iron Man 3
546     Minions
26      Captain America: Civil War
Name: title, dtype: object
```

2.0 Recommendation Engines

2.1 Initial Data Cleaning

```
In [5]: movies_credits.drop('title', axis = 1, inplace = True)

        movies_credits.head()
```

Out[5]:

	movie_id	cast	crew
0	19995	[{"cast_id": 242, "character": "Jake Sully", "...	[{"credit_id": "52fe48009251416c750aca23", "de...
1	285	[{"cast_id": 4, "character": "Captain Jack Spa...	[{"credit_id": "52fe4232c3a36847f800b579", "de...
2	206647	[{"cast_id": 1, "character": "James Bond", "cr...	[{"credit_id": "54805967c3a36829b5002c41", "de...
3	49026	[{"cast_id": 2, "character": "Bruce Wayne / Ba...	[{"credit_id": "52fe4781c3a36847f81398c3", "de...
4	49529	[{"cast_id": 5, "character": "John Carter", "c...	[{"credit_id": "52fe479ac3a36847f813eaa3", "de...

```
In [6]: movies_credits.columns = ['id', 'cast', 'crew']
movies = movies.merge(movies_credits, on='id')

movies.head()
```

Out[6]:

	budget	genres	homepage	id	keywords	original_language
0	237000000	[[{"id": 28, "name": "Action"}, {"id": 12, "nam...	http://www.avatarmovie.com/	19995	[[{"id": 1463, "name": "culture clash"}, {"id":...	en
1	300000000	[[{"id": 12, "name": "Adventure"}, {"id": 14, "...	http://disney.go.com /disneypictures/pirates/	285	[[{"id": 270, "name": "ocean"}, {"id": 726, "na...	en
2	245000000	[[{"id": 28, "name": "Action"}, {"id": 12, "nam...	http://www.sonypictures.com /movies/spectre/	206647	[[{"id": 470, "name": "spy"}, {"id": 818, "name...	en
3	250000000	[[{"id": 28, "name": "Action"}, {"id": 80, "nam...	http://www.thedarkknighttrises.com/	49026	[[{"id": 849, "name": "dc comics"}, {"id": 853,...	en
4	260000000	[[{"id": 28, "name": "Action"}, {"id": 12, "nam...	http://movies.disney.com/john- carter	49529	[[{"id": 818, "name": "based on novel"}, {"id":....	en

5 rows × 22 columns

2.2 Demographic Filtering - Trending now for All Users

```
In [7]: C = movies['vote_average'].mean()

C
```

Out[7]: 6.092171559442011

```
In [8]: m = movies['vote_count'].quantile(0.8)

m
```

```
Out[8]: 957.60000000000004
```

```
In [9]: qualified_movies_demographic = movies.copy().loc[movies['vote_count'] >= m]

qualified_movies_demographic.shape
```

```
Out[9]: (961, 22)
```

```
In [10]: def weighted_rating_score(data, m=m, C=C):
          v = data['vote_count']
          R = data['vote_average']
          # weighted rating score calculation formula
          return (v/(v+m) * R) + (m/(m+v) * C)
```

```
In [11]: # calculate weighted rating score for each movie we have

qualified_movies_demographic['weighted_score'] = qualified_movies_demographic.appl
y(weighted_rating_score, axis=1)
```

```
In [12]: # descending sort movies based on the weighted rating score

qualified_movies_demographic = qualified_movies_demographic.sort_values('weighted_
score', ascending=False)
```

```
In [13]: # print out the top 20 movies based on the weighted_score

qualified_movies_demographic[['title', 'vote_count', 'vote_average', 'weighted_score']].head(20)
```

Out[13]:

	title	vote_count	vote_average	weighted_score
1881	The Shawshank Redemption	8205	8.5	8.248353
662	Fight Club	9413	8.3	8.096134
3337	The Godfather	5893	8.4	8.077404
3232	Pulp Fiction	8428	8.3	8.074738
65	The Dark Knight	12002	8.2	8.044250
809	Forrest Gump	7927	8.2	7.972814
96	Inception	13752	8.1	7.969290
95	Interstellar	10867	8.1	7.937399
1990	The Empire Strikes Back	5879	8.2	7.904757
1818	Schindler's List	4329	8.3	7.900080
3865	Whiplash	4254	8.3	7.894325
329	The Lord of the Rings: The Return of the King	8064	8.1	7.886879
2294	Spirited Away	3840	8.3	7.859318
2912	Star Wars	6624	8.1	7.846400
1553	Se7en	5765	8.1	7.813995
262	The Lord of the Rings: The Fellowship of the Ring	8705	8.0	7.810927
2731	The Godfather: Part II	3338	8.3	7.807818
690	The Green Mile	4048	8.2	7.796760
330	The Lord of the Rings: The Two Towers	7487	8.0	7.783656
77	Inside Out	6560	8.0	7.756979

2.3 Content Based Filtering - using Movie Overviews

```
In [14]: # use scikit-learn's TfidfVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer

# remove english stop words like 'the', 'a' and define a TF-IDF Vectorizer Object.
tfidf = TfidfVectorizer(stop_words='english')

# change NaN into an empty string
movies['overview'] = movies['overview'].fillna('')

# build the TF-IDF matrix
tfidf_matrix = tfidf.fit_transform(movies['overview'])
```

```
In [15]: # import linear_kernel
from sklearn.metrics.pairwise import linear_kernel

# build the cosine similarity matrix
cosine_sim_overviews = linear_kernel(tfidf_matrix, tfidf_matrix)
```

```
In [16]: # build a reverse map of movie titles and indices
indices = pd.Series(movies.index, index=movies['title']).drop_duplicates()
```

```
In [17]: # recommend similar movies to users given a movie title input

def get_recommendation_movies(title, cosine_sim):
    # get the movie index
    idx = indices[title]

    # calculate the pairwise similarity scores with that movie
    sim_scores = list(enumerate(cosine_sim[idx]))

    # sort the movies based on the similarity scores
    sim_scores = sorted(sim_scores, key=lambda x: x[1], reverse=True)

    # get the 10 most similar movies scores
    sim_scores = sim_scores[1:11]

    # get movie indices
    movie_indices = [i[0] for i in sim_scores]

    # return the recommended similar movies
    return movies['title'].iloc[movie_indices]
```

```
In [18]: get_recommendation_movies('The Dark Knight Rises', cosine_sim_overviews)
```

```
Out[18]: 65                The Dark Knight
299                Batman Forever
428                Batman Returns
1359                Batman
3854    Batman: The Dark Knight Returns, Part 2
119                Batman Begins
2507                Slow Burn
9        Batman v Superman: Dawn of Justice
1181                JFK
210                Batman & Robin
Name: title, dtype: object
```

```
In [19]: get_recommendation_movies('The Avengers', cosine_sim_overviews)
```

```
Out[19]: 7                Avengers: Age of Ultron
3144                Plastic
1715                Timecop
4124                This Thing of Ours
3311                Thank You for Smoking
3033                The Corruptor
588    Wall Street: Money Never Sleeps
2136    Team America: World Police
1468                The Fountain
1286                Snowpiercer
Name: title, dtype: object
```

2.4 Content Based Filtering - using Movie Metadata ('genres', 'keywords', 'cast', 'crew', 'production_companies')

```
In [20]: # import function to parse the stringified features into python objects

from ast import literal_eval
```

```
In [21]: features = ['genres', 'keywords', 'cast', 'crew', 'production_companies']

for feature in features:
    movies[feature] = movies[feature].apply(literal_eval)
```

```
In [22]: movies.head(3)
```

Out[22]:

	budget	genres	homepage	id	keywords	original_language	original_title
0	237000000	[{'id': 28, 'name': 'Action'}, {'id': 12, 'name': 'Adventure'}, {'id': 14, 'name': 'Fantasy'}]	http://www.avatarmovie.com/	19995	[{'id': 1463, 'name': 'culture clash'}, {'id': 1464, 'name': 'culture clash'}]	en	Avatar
1	300000000	[{'id': 12, 'name': 'Adventure'}, {'id': 14, 'name': 'Fantasy'}]	http://disney.go.com/disneypictures/pirates/	285	[{'id': 270, 'name': 'ocean'}, {'id': 726, 'name': 'pirates'}]	en	Pirates of the Caribbean: At World's End
2	245000000	[{'id': 28, 'name': 'Action'}, {'id': 12, 'name': 'Adventure'}, {'id': 14, 'name': 'Fantasy'}]	http://www.sonypictures.com/movies/spectre/	206647	[{'id': 470, 'name': 'spy'}, {'id': 818, 'name': 'thriller'}]	en	Spectre

3 rows × 8 columns

```
In [23]: # helper function: extract director name

def get_director(crew_data):
    for i in crew_data:
        if i['job'] == 'Director':
            return i['name']
    return np.nan
```

In [24]: *# helper function: transfer into list data objects*

```
def get_list(data):
    if isinstance(data, list):
        names = [i['name'] for i in data]
        # return only the first three items
        if len(names) > 3:
            names = names[:3]
        return names

    return []
```

In [25]: *# use helper functions to parse features*

```
movies['director'] = movies['crew'].apply(get_director)

features = ['genres', 'keywords', 'cast', 'production_companies']
for feature in features:
    movies[feature] = movies[feature].apply(get_list)
```

In [26]: *# print the new features of the first 3 films*

```
movies[['title', 'genres', 'keywords', 'cast', 'production_companies', 'director']]
.head(3)
```

Out[26]:

	title	genres	keywords	cast	production_companies	director
0	Avatar	[Action, Adventure, Fantasy]	[culture clash, future, space war]	[Sam Worthington, Zoe Saldana, Sigourney Weaver]	[Ingenious Film Partners, Twentieth Century Fo...]	James Cameron
1	Pirates of the Caribbean: At World's End	[Adventure, Fantasy, Action]	[ocean, drug abuse, exotic island]	[Johnny Depp, Orlando Bloom, Keira Knightley]	[Walt Disney Pictures, Jerry Bruckheimer Films...]	Gore Verbinski
2	Spectre	[Action, Adventure, Crime]	[spy, based on novel, secret agent]	[Daniel Craig, Christoph Waltz, Léa Seydoux]	[Columbia Pictures, Danjaq, B24]	Sam Mendes

In [27]: *# convert all strings into lower case and
strip the spaces between names so as to be specific*

```
def clean_data(data):
    if isinstance(data, list):
        return [str.lower(i.replace(" ", "")) for i in data]
    else:
        if isinstance(data, str):
            return str.lower(data.replace(" ", ""))
        else:
            return ''
```

In [28]: features = ['genres', 'keywords', 'cast', 'production_companies', 'director']

```
for feature in features:
    movies[feature] = movies[feature].apply(clean_data)
```



```
In [29]: def create_data_soup(x):
         return ' '.join(x['keywords']) + ' ' + ' '.join(x['cast']) + ' ' + x['director'] + ' ' + ' '.join(x['genres']) + ' ' + ' '.join(x['production_companies'])
```

```
In [30]: movies['data_soup'] = movies.apply(create_data_soup, axis=1)
```

```
In [31]: # import CountVectorizer, create count matrix
         from sklearn.feature_extraction.text import CountVectorizer

         count = CountVectorizer(stop_words='english')
         count_matrix = count.fit_transform(movies['data_soup'])
```

```
In [32]: # compute the Cosine Similarity matrix
         from sklearn.metrics.pairwise import cosine_similarity

         cosine_sim_metadata = cosine_similarity(count_matrix, count_matrix)
```

```
In [33]: # reset the indices of our main DataFrame, construct reverse mapping

         movies = movies.reset_index()
         indices = pd.Series(movies.index, index=movies['title'])
```

```
In [34]: get_recommendation_movies('The Dark Knight Rises', cosine_sim_metadata)
```

```
Out[34]: 65          The Dark Knight
         119          Batman Begins
         14          Man of Steel
         1196         The Prestige
         4638  Amidst the Devil's Wings
         10          Superman Returns
         1035          Jonah Hex
         299          Batman Forever
         303          Catwoman
         747          Gangster Squad
         Name: title, dtype: object
```

```
In [35]: get_recommendation_movies('The Avengers', cosine_sim_metadata)
```

```
Out[35]: 7          Avengers: Age of Ultron
         26          Captain America: Civil War
         79          Iron Man 2
         169  Captain America: The First Avenger
         85  Captain America: The Winter Soldier
         174          The Incredible Hulk
         31          Iron Man 3
         68          Iron Man
         182          Ant-Man
         94          Guardians of the Galaxy
         Name: title, dtype: object
```

2.5 Content Based Filtering - using Movie Metadata ('genres', 'keywords', 'cast', 'crew', 'production_companies') and voting scores

```
In [39]: def improved_recommendations_metadata_votes(title, cosine_sim):
    idx = indices[title]
    sim_scores = list(enumerate(cosine_sim[idx]))
    sim_scores = sorted(sim_scores, key=lambda x: x[1], reverse=True)
    sim_scores = sim_scores[1:36]
    movie_indices = [i[0] for i in sim_scores]

    movies_improved_recomm = movies.copy().iloc[movie_indices][['title', 'vote_count', 'vote_average']]
    vote_counts_improved = movies_improved_recomm[movies_improved_recomm['vote_count'].notnull()][['vote_count']].astype('int')
    vote_averages_improved = movies_improved_recomm[movies_improved_recomm['vote_average'].notnull()][['vote_average']].astype('int')
    C_improved = vote_counts_improved.mean()
    m_improved = vote_averages_improved.quantile(0.60)

    qualified_improve = movies_improved_recomm[(movies_improved_recomm['vote_count'] >= m_improved) & (movies_improved_recomm['vote_count'].notnull()) & (movies_improved_recomm['vote_average'].notnull())
    qualified_improve['wr'] = qualified_improve.apply(weighted_rating_score, axis=1)
    qualified_improve = qualified_improve.sort_values('wr', ascending=False).head(10)
    return qualified_improve
```

```
In [40]: improved_recommendations_metadata_votes('The Dark Knight Rises', cosine_sim_metadata)
```

/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:15: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

```
from ipykernel import kernelapp as app
```

Out[40]:

	title	vote_count	vote_average	wr
65	The Dark Knight	12002	8.2	8.044250
96	Inception	13752	8.1	7.969290
95	Interstellar	10867	8.1	7.937399
1196	The Prestige	4391	8.0	7.658427
119	Batman Begins	7359	7.5	7.337898
1663	Once Upon a Time in America	1069	8.2	7.204018
1052	Training Day	1634	7.3	6.853706
3854	Batman: The Dark Knight Returns, Part 2	419	7.9	6.642426
14	Man of Steel	6359	6.5	6.446623
1456	Bound by Honor	115	7.7	6.264557

```
In [41]: improved_recommendations_metadata_votes('The Avengers', cosine_sim_metadata)
```

```
/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:15: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy
from ipykernel import kernelapp as app
```

```
Out[41]:
```

	title	vote_count	vote_average	wr
94	Guardians of the Galaxy	9742	7.9	7.738202
85	Captain America: The Winter Soldier	5764	7.6	7.385186
68	Iron Man	8776	7.4	7.271335
158	Star Trek	4518	7.4	7.171280
47	Star Trek Into Darkness	4418	7.4	7.167026
7	Avengers: Age of Ultron	6767	7.3	7.150268
26	Captain America: Civil War	7241	7.1	6.982285
182	Ant-Man	5880	7.0	6.872859
1294	Serenity	1264	7.4	6.836273
31	Iron Man 3	8806	6.8	6.730577